

CS2102 Cheatsheet
Selection – Select all tuples from relation that satisfy condition iff evaluates to **true**

$\sigma_{\text{manager}='Judy'}(\text{Projects})$

Selection Conditions

attribute op constant $\sigma_{\text{start_year}=2020}(\text{Projects})$
attribute, op attribute₂ $\sigma_{\text{start_year}=\text{end_year}}(\text{Projects})$
expr₁ ∧ expr₂ $\sigma_{\text{start_year}=2020 \wedge \text{manager}='Judy'}(\text{Projects})$
expr₁ ∨ expr₂ $\sigma_{\text{start_year}=2020 \vee \text{manager}='Judy'}(\text{Projects})$
¬ expr $\sigma_{\neg(\text{start_year}=2020)}(\text{Projects})$

→ **Comparison operation with null = unknown**
→ **Arithmetic operation with null = null**
Projection

$\pi_{\text{pname}, \text{ename}}(\text{Teams})$

→ Order matters, **duplicates removed**

Renaming $\rho_{\text{pname} \leftarrow \text{ename}, \text{title} \leftarrow \text{pname}}(\text{Teams})$

$B_i \leftarrow A_i, \dots, B_k \leftarrow A_k$ final ← original

Set Operators

Union Intersection Set Difference

$R \cup S \quad R \cap S \quad R - S$

→ R and S must be union compatible, i.e. they have the same number of column attributes and must be compatible type (no need same name)

Cross Product X

→ A X B returns {J,K,Y,Z} for columns J, K belonging to A and Y, Z belonging to B

Inner Joins $R \bowtie_{\theta} S = \sigma_{\theta}(R \times S)$

→ Join based on condition θ specified

Projects $\pi_{\text{manager}=\text{name}} \text{Managers}$

→ Manager belongs to Projects, name belongs to Managers. Output has **both** manager & name

→ Above is also an equi join

→ Inner Joins allow for arbitrary comparison operators (e.g., =, <>, <, <=, >=, >)

Equi Joins

Special case of inner join – only for ‘=’ operator

Natural Join

→ Over **ALL** attributes R and S have in common, **output contains common attributes only once**

$R \bowtie S = \pi_{\ell}(R \bowtie_{\text{c}} \rho_{b_1 \leftarrow a_1, \dots, b_k \leftarrow a_k}(S))$

Outer Join

- Includes dangling tuples + inner join

Perform inner join $M = R \bowtie_{\theta} S$

To M, add dangling tuples to result of

R in case of a **left outer join** \bowtie_{θ}

S in case of a **right outer join** \bowtie_{θ}

R and S in case of a **full outer join** \bowtie_{θ}

Employees $\bowtie_{\text{manager}=\text{ename}} (\pi_{\text{ename}}(\text{Projects}))$

name	age	role	ename
Jack	40	dev	Jack
Bill	45	dev	Bill
Marie	36	hr	null
Bernie	19	null	null

} inner join result

} dangling tuples with null padding

→ Basically match all employees based on condition to projects table, and include those who ended up not matching to anything
→ Full outer joins have dangling tuples padded with null for both tables R and S.

Natural Outer Join

Outer join that outputs common attributes of R

and S once, similar to


Natural Join

Natural left outer join $R \bowtie_{\text{L}} S$

Natural right outer join $R \bowtie_{\text{R}} S$

Natural full outer join $R \bowtie_{\text{F}} S$

Relational Operator Trees




ER MODEL - Attributes

Specific info describing an entity – represented by oval in ER diagrams. 4 types:

Key attribute (uniquely identifies entity), **Composite attribute** (comprised of multiple attributes, oval branched to other ovals), **Multivalued attribute** (consists of more than one value, double-lined oval),

Derived attribute (derived from other attributes, dashed oval)

Relationship / Relationship Sets (Roles)




→ Descriptor of entity set's participation in a rs, explicit role label only common in ambiguity or if same entity set participates in relationship > 1 time

→ Rectangle = object, Diamond = relationship

n-ary Relationship Set

n = degree of relationship set, involves n entity roles (> 3 are very rare)
n = 2 → binary relationship set
n = 3 → ternary relationship set




Cardinality Constraints


- Many to Many (m:n), Many to one (m:1), One to one (1:1)

Participation Constraints

- **Partial participation constraint (default)**
 - Participation of an entity in a relationship is not mandatory
 - Example: A user made 0 or more bookings



E is a weak entity set with identifying owner 'E' and identifying relationship set R.



Weak Entity Sets



CREATE TABLE Flights (...)
CREATE TABLE FlightInstances (...)

→ Do not have own key, can only be uniquely identified by considering pk of owner entity

→ Weak entity's existence depends on existence of owner entity

→ Many to one relationship from weak entity set to owner entity set

→ **Partial Key** – Set of attributes of weak entity that uniquely identifies a weak entity for a given owner entity

Modelling Cardinality Constraints → Schema

M:N → Create table with pk as all the attributes of the entity sets participating, have fks pointing back to the entity sets

1:M → Create table with the entity participating once as pk, have fk pointing to both entities

1:M → Combine relationship set and entity set participating once, point FK to parent

1:1 → Put just them all in the same table, or if you want to split them, have FKs pointing to each other


ISA Hierarchies - Overlap Constraint

→ Can a superclass entity belong to multiple subclasses (can be both vs either student/staff)


ISA Hierarchies - Covering Constraint

→ Does superclass entity have to belong to a subclass (Every student must be a grad / undergrad vs not every person is a student)

ES




Example



→ Use ON DELETE CASCADES / FKs etc to model

Aggregation

→ Treat relationships as higher level entities



Functional Dependencies

Armstrong Axioms

1. Axiom of Reflexivity (Set of attributes → subset of attributes)

- {NRIC, Name} → {NRIC}
- {StudentID, Name, Age} → {Name, Age}

2. Axiom of Augmentation (If A → B then AC → BC)

- **Example:** if {NRIC} → {Name} then
 - {NRIC, Age} → {Name, Age}
 - {NRIC, Salary, Weight} → {Name, Salary, Weight}
 - {NRIC, Address, Postal} → {Name, Address, Postal}

3. Axiom of Transitivity

- If {A} → {B} and {B} → {C}
- Then {A} → {C}

Example:

- if {NRIC} → {Address}
- and {Address} → {Postal}
- then {NRIC} → {Postal}

Extended Armstrong's Axioms


4. Rule of Decomposition

- If {A} → {BC}
- Then {A} → {B} and {A} → {C}

5. Rule of Union

- If {A} → {B} and {A} → {C}
- Then {A} → {BC}

→ Can use circuit diagram to compute closures



Boyce-Codd Normal Form (BCNF)

- Non-trivial means when your FD/closure doesn't just point to itself (e.g. A->A)

- Decomposed means the RHS only has 1 attribute (e.g. AB->C, A->D)

- BCNF definition: Every non-trivial and decomposed FD has a superkey as its LHS

BCNF Decomposition

- Compute closure for all possible combinations of attributes, find one closure where RHS is not trivial and not all attributes

- If schema is (ABCD) and FD violation is A->ABC, split tables into ABC and AD

→ Take the 3 violating attributes put in one table, the other table has the LHS attribute + the attribute(s) that didn't appear

→ Keep splitting until every single table has either 2 attributes or is in BCNF

→ If the new table you split don't have relevant FDs, just use it as if its there and remove it after

iii. Check if R1(A, B) and R2(A, C, D, E) are in BCNF

YES! ???
Do we even know what are the FDs that holds on R2?

{A} → {B}? But no B {BC} → {D}? Also no B

a. **Compute closure of each subset of attributes**

b. **Remove attributes not in the current table** (projection)

- Now we can find the violation

{A}⁺ = {A} {C}⁺ = {C} {D}⁺ = {D} {E}⁺ = {E}
{AC}⁺ = {ABCD} {AD}⁺ = {ABD} {AE}⁺ = {AE}
{CD}⁺ = {CD} {CE}⁺ = {CE} {DE}⁺ = {DE}
{ACD}⁺ = {ABCD} {ACE}⁺ = {ABCD} {ADE}⁺ = {ABCD} {CDE}⁺ = {CDE}

→ From the image, we split ACDE into ACD and ACE

Properties of BCNF

- **Good properties**
 - No update or deletion or insertion anomalies
 - Small redundancies (very hard to completely remove redundancies)
 - The original table can always be reconstructed from the decomposed tables
- **Bad properties** (non dependency preserving decomposition)
 - Dependencies may not be preserved
 - Topic for next week lecture

3NF

→ BCNF does not preserve dependencies

→ A table is in 3NF is every non-trivial and decomposed FD either has LHS as superkey or RHS is a prime attribute

Find the keys of R:

1. Look at the RHS, if the attribute not there then it will definitely be part of key

2. Try use FDs, start from the smallest ones, get the closures, get minimal ones that give us closure with all attributes (if we can find 2 attributes to give us the key then 2 is the key size, anything else is superkey)

3. All attributes in key are prime attributes

(a) Find all the keys of R. You do not have to show your step. Write the answer in the following format:

{A, B, C, D}, {B, C, D, E}, {C, D, E, A}

Solution: All the keys are {A, C, E}, {A, C, D}

Minimal Basis:

1. See if we can form all the FDs that are defined in the question based on the few that we picked
2. Try finding as little as possible

(c) Find a minimal basis of the set of functional dependencies. You do not have to show your step. Write the answer in the following format:

{A, B, C} → {C, D, E}, {B, C, D} → {D, E, A}, {C, D, E} → {E, A, B}

Solution: One possible minimal basis is:

{{B, D} → {E}, {C, D} → {B}, {C, E} → {D}}

Lossless join and dependency preserving 3NF decomposition of R

1. Combine the FD we got from the minimal basis. Check the keys we calculated, there must be **at least one key** from the schema and FDs in the 3NF decomposition, so if its not present, add it in as a new table.
2. Remember when combine must create **canonical cover**

(d) Using the minimal basis you computed in part (c), find a lossless-join and dependency-preserving 3NF decomposition of R. Note that you have to first ensure that your answer to part (c) is a minimal basis first. Write the answer in the same format as part (b).

Solution: Using the previous minimal basis, we arrive at the following decomposition:
{R1(B, D, E), R2(B, C, D), R3(C, D, E)}

The keys are not in the decomposition, so we add the key R4(A, C, E). And viola, we have exactly the same decomposition as our BCNF decomposition.

Lossless join and dependency preserving BCNF decomposition of R

→ Typically this is by luck, may not always get one. We get it by checking our 3NF decomposition, make sure every table is in BCNF.

Deriving preserved FDs from split tables

- Compute closure on each table w.r.t original fds!!! → get the new fds that are captured by the new tables. Use them to check if they can derive the original FDs.

(a) Consider the following ER diagram:




Figure 1: ER Diagram

→ Qn: Find min set of attr that can uniquely identify all other attr

Solution: A = {a1, s1, c1}

Due to the key constraint from A to R, we actually have {a1} → {r1}.

→ Notice that we have a participation and cardinality constraint from A to R, so there is only one unique a1, which means a1 → r1, b1, so we don't need b1 to uniquely identify the other attributes in R

→ **double line** – need both keys. If have arrow and double line then **only need that key**. If question says ignore NULL value then I think arrow itself can imply FD alrd

A lossless join decomposition is a decomposition of a relation into relations. such that a natural join of the two smaller relations yields back the original relation

To check for lossless join decomposition using FD set, e.g R(ABCD) splits into R1(ABC), R2(AD), R3(CD). For R1 and R2, common attribute is A. Check if A → ABC or A → AD can be found using the FDs from the original table ABCD. If can compute with the closure → it is lossless

SQL Data Types

boolean	logical Boolean (true/false)
integer	signed four-byte integer
float8	double precision floating-point number (8 bytes)
numeric [(p,s)]	exact numeric of selectable precision
char(n)	fixed-length character string
varchar(n)	variable-length character string
text	variable-length character string
date	calendar date (year, month, day)
timestamp	date and time

Numeric(10,3) – 10 digits, 3 d.p. If u put 1d.p pgsql will auto pad 0s until 10dp

Inserting, Deleting, Update-ing Data:

INSERT INTO Employees (id, name) VALUES (102, 'Judy'), (103, 'Max');

-- Delete selected tuples
DELETE FROM Employees
WHERE role = 'dev';

```
UPDATE Sells s
SET price = CASE
    WHEN r.area = 'Central' THEN price + 3
    WHEN r.area = 'EAST' THEN price + 2
    ELSE price + 1 END
FROM Restaurants r
where s.manager = r.name;
```

SQL NULL Values

Use IS NULL to check if SQL value is null. For non null, can use x IS NOT NULL or NOT (x IS NULL)

IS DISTINCT FROM

Equivalent to x <> y. For x is null and y is null, x IS DISTINCT FROM y returns false, for x null and y non-null, x IS DISTINCT FROM y returns true

NOT NULL Constraint

```
CREATE TABLE Employees (
    id VARCHAR(50) CONSTRAINT nn_id NOT NULL,
    name VARCHAR(50) CONSTRAINT nn_name NOT NULL,
    age INTEGER,
    role VARCHAR(50),
);
```

UNIQUE Constraint

can be used across multiple columns in table constraint

unnamed column constraint

```
CREATE TABLE Employees (
    id INTEGER UNIQUE,
    name VARCHAR(50),
    age INTEGER,
    role VARCHAR(50)
);
```

named column constraint

```
CREATE TABLE Employees (
    id INTEGER CONSTRAINT u_id UNIQUE,
    name VARCHAR(50),
    age INTEGER,
    role VARCHAR(50)
);
```

unnamed table constraint

```
CREATE TABLE Employees (
    id INTEGER,
    name VARCHAR(50),
    age INTEGER,
    role VARCHAR(50),
    UNIQUE (id)
);
```

named table constraint

```
CREATE TABLE Employees (
    id INTEGER,
    name VARCHAR(50),
    age INTEGER,
    role VARCHAR(50),
    CONSTRAINT u_id UNIQUE (id)
);
```

For unique to be false:

"(t₁.eid <> t₂.eid) or (t₁.pname <> t₂.pname)" evaluates to **false**
i.e. both eid and pname for 2 tuples exactly same

Note

- Column and table constraints can be combined, even w.r.t to same column
 - All can be column/table constraint except NULL
- Foreign Key Constraints**
- Reference pkey of another table

```
CREATE TABLE Teams (
    eid INTEGER,
    pname VARCHAR(100),
    house INTEGER,
    PRIMARY KEY (eid, pname),
    FOREIGN KEY (eid) REFERENCES Employees (id) ON DELETE NO ACTION ON UPDATE CASCADE,
    FOREIGN KEY (pname) REFERENCES Projects (name) ON DELETE SET NULL ON UPDATE CASCADE
);
```

→ Might not need to specify id and name in employees and projects, DBMS know which is pk
→ Either primary key in referenced relation OR NULL value
→ Deletion/Update of referenced tuple can lead to error thrown/cascade etc

FK Violation Actions (ON DELETE/UPDATE)

NO ACTION

rejects delete/update if it violates constraint (default value)

RESTRICT

similar to "no action" except that check of constraint cannot be deferred (deferred constraints are discussed in a bit)

CASCADE

propagates delete/update to referencing tuples

SET DEFAULT

updates foreign keys of referencing tuples to some default value (important: default value must be a primary key in the referenced table)

SET NULL

updates foreign keys of referencing tuples to null (important: corresponding column must allow to contain null values)

CREATE TABLE Teams (
 eid INTEGER,
 pname VARCHAR(100) DEFAULT 'FastCash',
 house INTEGER,
 FOREIGN KEY (eid) REFERENCES Employees (id) ON UPDATE CASCADE,
 FOREIGN KEY (pname) REFERENCES Projects (name) ON UPDATE CASCADE ON DELETE SET NULL
);

optional since it's the default action

FK is optional

→ Special example: when project is deleted, set null, but default fastcash so final value is fastcash
→ ON DELETE CASCADE can have bad conseq, and affect performance

CHECK Constraints – CHECK(Condition)

→ can make condition as complex as needed

```
CREATE TABLE Projects (
    name VARCHAR(50) PRIMARY KEY,
    start_year INTEGER,
    end_year INTEGER,
    -- CHECK (start_year <= end_year),
    CONSTRAINT valid_lifetime CHECK (start_year <= end_year)
);
```

CREATE ASSERTIONS

Not really used in practice, but triggers are.

DEFERRABLE CONSTRAINTS

- Constraint check deferred to end of tranxn

Deferrable Constraints – Motivation

→ NOT DEFERRABLE, check constraint after each statement in a transaction
CONSTRAINT manager_fkey FOREIGN KEY (manager) REFERENCES Employees (id) NOT DEFERRABLE – default value (optional), check if constraint is immediate and cannot be changed
→ DEFERRABLE INITIALLY DEFERRED, only check constraint after transaction ends
CONSTRAINT manager_fkey FOREIGN KEY (manager) REFERENCES Employees (id) DEFERRABLE INITIALLY DEFERRED – check of constraint deferred by default
→ DEFERRABLE INITIALLY IMMEDIATE – DB checks if violation occurs after every statement, but we can switch it off on demand, unlike NOT DEFERRABLE, can never be switched off

```
INSERT INTO Employees VALUES (101, 'Sarah', 102), (102, 'Judy', 101), (103, 'Max', 102);

BEGIN;
SET CONSTRAINT manager_fkey DEFERRED; -- Set check of constraint from "immediate" to "def"
DELETE FROM Employees WHERE id = 102; -- Judy got fired → constraint violated but not checked
UPDATE Employees SET manager = 101 WHERE id = 103; -- Max gets a new manager → constraint re-evaluated
COMMIT;
```

+ No need to care abt order of statements within txn, allow for cyclic fk constraints, less checks higher performance
- Harder to debug, data definitn not unambiguous

Modify Single Column in DB

```
ALTER TABLE Projects ALTER COLUMN name TYPE VARCHAR(200); -- change data type to VARCHAR(200)
ALTER TABLE Projects ALTER COLUMN start_year SET DEFAULT 2021; -- set default value of column "start_year"
ALTER TABLE Projects ALTER COLUMN start_year DROP DEFAULT; -- drop default value of column "start_year"
```

Add/Drop Columns

```
ALTER TABLE Projects ADD COLUMN budget NUMERIC DEFAULT 0.0; -- add new column with a default value
ALTER TABLE Projects DROP COLUMN budget; -- drop column from table
```

Add/Drop Constraints

```
ALTER TABLE Teams ADD CONSTRAINT eid_fkey FOREIGN KEY (eid) REFERENCES Employees (id); -- add foreign key constraint
ALTER TABLE Teams DROP CONSTRAINT eid_fkey; -- drop foreign key constraint (name of constraint might be retrieved from metadata)
```

Dropping Tables

- DROP TABLE (IF EXISTS) Projects;
- With dependent objects (assume foreign key constraint Teams.pname → Projects.name)

```
DROP TABLE Projects; -- will throw an error because of foreign key constraint
DROPPING TABLES CASCADE; -- will delete table "Projects" and foreign key constraint (will not delete table "Teams")
```

→ Drop foreign key constraint in Teams, does not delete teams!
SELECT CLAUSE – DOESN'T ELIMINATE DUPS
SELECT [DISTINCT] target-list FROM relation-list
[WHERE conditions] [AND condition]
→ combine, process attributes, rename columns

Find the name and the all countries the GDP per capita in SGD rounded to the nearest dollar (on the database).

name	gdp_per_capita
Denmark	\$3.63955
France	\$3.59450
Kirgizstan	\$3.4952
Norway	\$3.9358
Singapore	\$3.13344
Sweden	\$3.4464
Turkmenistan	\$3.24366
United Arab Emirates	\$3.15110

"g" concatenates strings

"AB" is optional

Convert from USD to SGD

→ Use DISTINCT to enforce dup elimination

WHERE Clause

→ DO NOT use col = NULL as condition, returns unknown → no tuples returned (no err thrown); use column IS NULL / column IS NOT NULL
→ Pattern Matching
- "_" matches any single character
- "%" matches any seq of 0 / more characters
Find all cities that start with "Si" and end with "re".

```
SELECT name
FROM cities
WHERE name LIKE 'Si%re';
```

name
Singapore
Sierre
Sierra Madre

Set Operation Queries

→ Any SQL queries that yield union-compatible tables can do:
- Q1 UNION (ALL)
Q2=Q1 U Q2
- Q1 INTERSECT (ALL) Q2=Q1 ∩ Q2
- Q1 EXCEPT (ALL) Q2 = Q1 – Q2
→ Eliminates duplicates from result
→ Add ALL behind set operation if we do not want to eliminate tuples from result

Complex JOIN query example

Find all cities sorted by country (ascending from A to Z) and for each country with respect to the cities population in descending order.

```
SELECT n.name AS country, c.name AS city, c.population
FROM cities c, countries n
WHERE c.country_id = n.id
ORDER BY n.name ASC, c.population DESC;
```

Find all airports in European countries without a land border which cannot be reached by plane given the existing routes in the database.

```
SELECT t1.country, t1.city, t1.airport
FROM
    (SELECT n.name AS country, c.name AS city,
     a.name AS airport, a.code
     FROM borders b, countries n, cities c, airports a
     WHERE b.country_id1 = n.id1 AND n.id1 = c.country_id1 AND c.name = a.city AND b.country_id2 = n.id2 AND n.id2 = c.country_id2 AND n.continent = 'Europe') t1
LEFT OUTER JOIN
    routes r
ON t1.code = r.to_code
WHERE r.to_code IS NULL;
```

Annotations: "subquery", "table", "3 table join", "table selection by join condition", "All airports in European countries without a land border (13 tuples)", "as optional"

Subqueries

→ Must be enclosed in parentheses, table alias mandatory, column aliases optional
→ Subquery must return exactly 1 col (not row)
→ Expression compared to every subquery row
→ Subquery can be correlated to outer query
→ Can contain multiple nested subqueries
IN / NOT IN Subqueries
→ Can replace IN with inner joins, NOT IN with outer joins

```
SELECT *
FROM countries
WHERE continent IN ('Asia', 'Europe')
AND population BETWEEN 5000000 AND 6000000;
```

Find all names that refer to both a city and a country.

```
SELECT name
FROM countries
WHERE name IN (SELECT name
FROM cities);
```

outer query

inner query

ANY / SOME Subqueries

→ Just need to be true for 1 val in subquery
Find all countries with a population size smaller than any city called "London" (there are 4 cities called "London" on the database).

```
SELECT name, population
FROM countries
WHERE population < ANY (SELECT population
FROM cities
WHERE name = 'London');
```

145 tuples

ALL Subqueries

→ Must be true for ALL values in subquery
For each continent, find the country with the highest GDP.

SELECT name, continent, gdp FROM countries c1, SELECT gdp FROM countries c2 WHERE gdp >= ALL (SELECT gdp FROM countries c2 WHERE c2.continent = c1.continent);

BRUNY

name	continent	population
Singapore	Asia	5781728
France	Europe	66324761
Sweden	Europe	969390
Brunei	Asia	438263
Brunei	Asia	759252
Algeria	Africa	4017173
Senegal	Africa	2073440
Mexico	North America	320601

name	continent	population
Australia	Oceania	116000000000
Brazil	South America	306000000000
China	Asia	2110000000000
Egypt	Africa	111000000000
Germany	Europe	340000000000
United States	North America	1860000000000

Correlated Subqueries

→ Naming ambiguities: Use table aliases, but scoping rules go from inner to outer scope
→ Outer scope will not know about table alias declaration in inner scope
(NOT) EXISTS Subqueries
→ EXISTS – As long as one tuple match, NOT EXISTS – NO tuples can match
→ NOT EXISTS
subqueries should be correlated – good tool to filter tuples, similar to EXCEPT?
Scalar Subqueries
→ Subquery that returns 1 single value
→ DBMS looks at key constraints to know that subquery is scalar
→ Can use as SELECT list, WHERE, JOIN
Sorting – ORDER BY
→ ORDER BY column DESC / ASC
→ Can sort w.r.t multiple attributes

Find the all the countries for which there is no city in the database.

```
SELECT n.name
FROM countries n
WHERE NOT EXISTS (SELECT *
FROM cities c
WHERE c.country_id = n.id);
```

For all cities, find their names together with the names of the countries they are located in.

```
SELECT name AS city,
(SELECT name AS country
FROM countries n
WHERE n.id = c.country_id)
FROM cities c;
```

Find all cities sorted by country (ascending from A to Z) and for each country with respect to the cities population in descending order.

```
SELECT country, city, population
FROM cities
ORDER BY country ASC, population DESC;
```

country	city	population
Algeria	Alger	367506
Algeria	Kandahar	491023
Algeria	Heard	430303
Algeria	Traza	419465
Algeria	Dorra	124242
Algeria	Vire	120677
Zimbabwe	Chibwa	10283
Zimbabwe	Matipa	9886
Zimbabwe	Pumura	2145

Find all cities sorted by country (ascending from A to Z) and for each country with respect to the cities population in descending order.

```
SELECT n.name AS country, c.name AS city, c.population
FROM cities c, countries n
WHERE c.country_id = n.id
ORDER BY n.name ASC, c.population DESC;
```

The 3rd sorting criteria only affects result if 1st sorting criteria does not yield an unambiguous order already!

ROW Constructors

- Combine cols to compare > 1 col in subquery
Find all countries with a higher population > higher gdp than France & Germany

```
SELECT name, population, gdp
FROM countries
WHERE (population, gdp) > ANY (SELECT population, gdp
FROM countries
WHERE name IN ('Germany', 'France'));
```

LIMIT / OFFSET (Usually used with ORDER BY)

→ Limit k – Return first k
→ OFFSET i – consider ith tuple onwards
Find the "second" top-5 countries regarding their GDP per capita for all countries.

```
SELECT name, (gdp/population) AS gdp_per_capita
FROM countries
ORDER BY gdp_per_capita DESC
OFFSET 5
LIMIT 5;
```

AGGREGATION FUNCTIONS

- Let R, S be two relations with an attribute A
 - Let R be an empty relation
 - Let S be a non-empty relation with n tuples but only null values for A

Query	Result
SELECT MIN(A) FROM R;	null
SELECT MAX(A) FROM R;	null
SELECT AVG(A) FROM R;	null
SELECT SUM(A) FROM R;	null
SELECT COUNT(A) FROM R;	0
SELECT COUNT(*) FROM R;	0

Query	Result
SELECT MIN(A) FROM S;	null
SELECT MAX(A) FROM S;	null
SELECT AVG(A) FROM S;	null
SELECT SUM(A) FROM S;	null
SELECT COUNT(A) FROM S;	0
SELECT COUNT(*) FROM S;	n

[MIN, MAX, AVG, COUNT, SUM(attr)]
- All except COUNT will ignore null values
→ Return type dependent on column data type

- Examples
 - MIN(), MAX(): defined for all data types; return data type same as input data type
 - SUM(): defined for all numeric data types; SUM(INTEGER) – BIGINT, SUM(REAL) – REAL, ...
 - COUNT(): defined for all data types; COUNT(...) – BIGINT

GROUP BY

→ Always applied with aggregation

- Example:
 - Table R, with three attributes A, B, C

A	B	C
null	4	19
6	1	null
20	2	10
1	1	2
1	18	2
null	21	19
6	20	null

SELECT FROM R GROUP BY A, C;		
null	4	19
null	21	19
6	1	null
6	20	null
20	2	10
1	1	2
1	18	2

GROUP BY Restrictions to SELECT clause

- If column A_i of table R appears in the SELECT clause, one of the following conditions must hold:
 - A_i appears in the GROUP BY clause
 - A_i appears as input of an aggregation function in the SELECT clause
 - The primary key or a candidate-key of R appears in the GROUP BY clause
- HAVING CLAUSE (GROUP BY)**
- HAVING conditions
 - Conditions check for each group defined by GROUP BY clause
 - HAVING clause cannot be used without a GROUP BY clause
 - Conditions typically involve aggregate functions

Find all countries that have at least one city with a population size larger than the average population size of all European countries

```
SELECT n.name, n.continent
FROM countries n
WHERE c.country_id = n.id
GROUP BY n.name, n.continent
HAVING MAX(c.population) > (SELECT AVG(population)
FROM countries
WHERE continent = 'Europe');
```

name	continent
China	Asia
Mexico	North America
India	Asia
Egypt	Africa
Philippines	Asia
Russia	Europe
Thailand	Asia
Brazil	South America
South Korea	Asia
Indonesia	Asia
United States	North America

If column A_i of table R appears in the HAVING clause, one of the following conditions must hold:

- A_i appears in the GROUP BY clause
- A_i appears as input of an aggregation function in the HAVING clause
- The primary key or a candidate-key of R appears in the GROUP BY clause

→ Same logic as aggregation
CASE – Conditional Expressions

```
SELECT class, COUNT(*) AS city_count
FROM
    (SELECT name, CASE
        WHEN population > 10000000 THEN 'Super City'
        WHEN population > 5000000 THEN 'Mega City'
        WHEN population > 1000000 THEN 'Large City'
        WHEN population > 500000 THEN 'Medium City'
        ELSE 'Small City' END AS class
    FROM cities) t
GROUP BY class;
```

→ Transform column based on conditions
COALESCE – Conditional Expr. for NULL vals

- COALESCE(value1, value2, value3, ...)
 - Returns the first non-NULL value in the list of input arguments
 - Returns NULL if all values in the list of input arguments are NULL
- Example: SELECT COALESCE(null, 1, null, 2) → 1

Find the number of cities for each city type; consider cities with NULL for column "capital" as "other".

```
SELECT capital, COUNT(*) AS city_count
FROM
    (SELECT COALESCE(capital, 'other') AS capital
    FROM cities) t
GROUP BY capital;
```

capital	city_count
primary	202
other	17147
admin	3531
minor	3687

NULLIF – Conditional Expr. for NULL Values

```
NULLIF(v1, v2)
- Returns NULL if v1=v2; else return v1
```

```
SELECT MIN(NULLIF(gdp, 0)) AS min_gdp,
ROUND(AVG(NULLIF(gdp, 0))) AS avg_gdp
FROM countries;
```

min_gdp	avg_gdp
1500000	54323956636

Universal Quantification

→ Concept of a "FORALL" operator – person X has visited ALL countries (for e.g.)
→ Use "not exists (X except Y) or below:
"Find the names of all users that have visited all countries."

```
SELECT u.user_id, u.name
FROM users u, visited v
WHERE u.user_id = v.user_id
GROUP BY u.user_id
HAVING COUNT(*) = (SELECT COUNT(*) FROM countries);
```

user_id	name
103	Max
107	Emma

Recursive Queries

Find all airports that can be reached from SIN with 0..2 stops (iteration to max. 2 stops purely for performance reasons)

```
WITH RECURSIVE right_path AS (
    SELECT from_code, to_code, 0 AS stops
    FROM connections c
    WHERE from_code = 'SIN'
    UNION ALL
    SELECT c.from_code, c.to_code, p.stops+1
    FROM right_path p, connections c
    WHERE p.to_code = c.from_code
    AND p.stops <= 2
)
SELECT DISTINCT to_code, stops
FROM right_path
ORDER BY stops ASC;
```

to_code	stops
BKK	0
PHN	0
AKA	0
DLU	1
AMS	2
BKK	2
PER	2
ZYL	2

103 tuples

to_code	stops
BKK	0
PHN	0
AKA	0
DLU	1
AMS	2
BKK	2
PER	2
ZYL	2

927 tuples

to_code	stops
BKK	0
PHN	0
AKA	0
DLU	1
AMS	2
BKK	2
PER	2
ZYL	2

1,725 tuples

Deriving preserved FDs from split tables

- Compute closure on each table w.r.t original
fds!!! → get the new fds that are captured by
the new tables. Use them to check if they can